

▼ Importing Libraries

```
# import contractions library.
!pip install contractions

Collecting contractions
  Downloading contractions-0.1.73-py2.py3-none-any.whl (8.7 kB)
Collecting textsearch>=0.0.21 (from contractions)
  Downloading textsearch-0.0.24-py2.py3-none-any.whl (7.6 kB)
Collecting anyascii (from textsearch>=0.0.21->contractions)
  Downloading anyascii-0.3.2-py3-none-any.whl (289 kB)
----- 289.9/289.9 kB 5.7 MB/s eta 0:00:00
Collecting pyahocorasick (from textsearch>=0.0.21->contractions)
  Downloading pyahocorasick-2.0.0-cp310-cp310-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2010_x86_64.whl
----- 110.8/110.8 kB 14.8 MB/s eta 0:00:00
Installing collected packages: pyahocorasick, anyascii, textsearch, contractions
Successfully installed anyascii-0.3.2 contractions-0.1.73 pyahocorasick-2.0.0 textsearch-0.0.24
```



```
import pandas as pd
import numpy as np
# %load_ext nb_black

# library to suppress warnings or deprecation notes
import warnings

warnings.filterwarnings("ignore")

# import Regex, string and unicodedata.
import re, string, unicodedata

import contractions

# import BeautifulSoup.
from bs4 import BeautifulSoup

# import Natural Language Tool-Kit.
import nltk

# download Stopwords.
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
True

# import stopwords.
from nltk.corpus import stopwords

# import Tokenizer.
from nltk.tokenize import word_tokenize, sent_tokenize

# library to split data
from sklearn.model_selection import train_test_split, StratifiedKFold

# libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno

# import wordcloud
import wordcloud
from wordcloud import STOPWORDS
from wordcloud import WordCloud
```

```
# remove the limit for the number of displayed columns
pd.set_option("display.max_columns", None)

# set the limit for the number of displayed rows
pd.set_option("display.max_rows", 200)

# to get diferent metric scores
from sklearn.metrics import (
    recall_score,
    accuracy_score,
    confusion_matrix, classification_report,
    f1_score,
    precision_score,
    precision_recall_fscore_support
)

# import vectorizers
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# import rfc and cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score

# import word preprocessors
from nltk.tokenize import word_tokenize
from nltk.stem import LancasterStemmer, WordNetLemmatizer
```

Loading the dataset

```
df = pd.read_csv('/content/Tweets.csv')
```

```
df.head()
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	negativereas
0	570306133677760513	neutral	1.0000	N
1	570301130888122368	positive	0.3486	N
2	570301083672813571	neutral	0.6837	N
3	570301031407624196	negative	1.0000	Bad Fliq
4	570300817074462722	negative	1.0000	Can't

```
texts = [[word.lower() for word in text.split()] for text in df]
```

```
df.head()
```

```

      tweet_id  airline_sentiment  airline_sentiment_confidence  negativereason
0  570306133677760513          neutral          1.0000          N
4  570306133677760513          neutral          1.0000          N

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14640 entries, 0 to 14639
Data columns (total 15 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   tweet_id                             14640 non-null  int64
 1   airline_sentiment                     14640 non-null  object
 2   airline_sentiment_confidence          14640 non-null  float64
 3   negativereason                        9178 non-null   object
 4   negativereason_confidence             10522 non-null  float64
 5   airline                               14640 non-null  object
 6   airline_sentiment_gold                 40 non-null     object
 7   name                                  14640 non-null  object
 8   negativereason_gold                   32 non-null     object
 9   retweet_count                         14640 non-null  int64
10   text                                  14640 non-null  object
11   tweet_coord                           1019 non-null   object
12   tweet_created                         14640 non-null  object
13   tweet_location                        9907 non-null   object
14   user_timezone                         9820 non-null   object
dtypes: float64(2), int64(2), object(11)
memory usage: 1.7+ MB

```

Observations:

There are 15 columns in the dataset. Half of the columns have null values. Considering both dependent and independent variables not having any null values, we will not do any null value processing. Most columns in the dataset are of object type. airline_sentiment is our dependent / target variable. text column is our independent variable that we will use for analysis. All other columns will be dropped at a later stage.

```
df.isnull().sum()
```

```

tweet_id          0
airline_sentiment 0
airline_sentiment_confidence 0
negativereason    5462
negativereason_confidence 4118
airline           0
airline_sentiment_gold 14600
name              0
negativereason_gold 14608
retweet_count     0
text              0
tweet_coord       13621
tweet_created     0
tweet_location    4733
user_timezone     4820
dtype: int64

```

```
df.isnull().sum() / len(df) * 100
```

```


tweet_id          0.000000
airline_sentiment 0.000000
airline_sentiment_confidence 0.000000
negativereason    37.308743
negativereason_confidence 28.128415
airline           0.000000
airline_sentiment_gold 99.726776
name              0.000000
negativereason_gold 99.781421
retweet_count     0.000000
text              0.000000
tweet_coord       93.039617
tweet_created     0.000000
tweet_location    32.329235
user_timezone     32.923497
dtype: float64

```

Data Visualization

```
msno.matrix(df)
```

```
<Axes: >
```



```
#Visualization of missing value using heatmap
plt.figure(figsize=(12,7))
sns.heatmap(df.isnull(), cmap = "Blues")
plt.title("Missing values?", fontsize = 15)
plt.show()
```



Interestingly, the only non-null values of the `_gold` columns seems to be the same entries for the most part. Meanwhile, there is some but not total overlap between location and timezone in terms of missing values.



```
print("Percentage null or na values in df")
((df.isnull() | df.isna()).sum() * 100 / df.index.size).round(2)
```

```
Percentage null or na values in df
tweet_id                0.00
airline_sentiment       0.00
airline_sentiment_confidence 0.00
negativereason          37.31
negativereason_confidence 28.13
airline                 0.00
airline_sentiment_gold  99.73
name                   0.00
negativereason_gold     99.78
retweet_count           0.00
text                   0.00
tweet_coord            93.04
tweet_created           0.00
tweet_location          32.33
user_timezone           32.92
dtype: float64
```

`airline_sentiment_gold`, `negativereason_gold` have more than 99% missing data And `tweet_coord` have nearly 93% missing data. It will be better to delete these columns as they will not provide any constructive information

```
del df["tweet_coord"]
```

```
del df["airline_sentiment_gold"]
del df["negativereason_gold"]
```

```
df.head()
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	negativereason
0	570306133677760513	neutral	1.0000	N
1	570301130888122368	positive	0.3486	N
2	570301083672813571	neutral	0.6837	N
3	570301031407624196	negative	1.0000	Bad Flight
4	570300817074462722	negative	1.0000	Can't find

```
freq = df.groupby("negativereason").size()
```

we can't fill it will affect in bad way for example we have positive review and we fill the values with mode that means with Customer Service Issue it is mismatch and can be affect on train model so we keep the data as it is.

```
# Checking duplicates
df.duplicated().sum()
```

```
39
```

```
# Dropping duplicates
df.drop_duplicates(inplace = True)

df.duplicated().sum()

0

df.sample(n = 10)
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	negative
2415	569218838190907393	positive	0.6487	
1069	569937143314849792	negative	1.0000	Di L
55	569996412286582784	negative	0.6939	Flight I P
293	568840560347373569	positive	1.0000	
983	569976131748823040	negative	1.0000	La
8786	567848171155496960	negative	0.6551	C
6325	568055480226066434	negative	1.0000	Ci Servic
1223	569878139003740163	negative	1.0000	Ci Servic
14601	569592830307508224	negative	1.0000	La
7354	569648229295329280	positive	1.0000	

```
df.describe().T
```

	count	mean	std	min
tweet_id	14601.0	5.692156e+17	7.782706e+14	5.675883e+17
airline_sentiment_confidence	14601.0	8.999022e-01	1.629654e-01	3.350000e-01
negativereason_confidence	10501.0	6.375749e-01	3.303735e-01	0.000000e+00
retweet_count	14601.0	8.280255e-02	7.467231e-01	0.000000e+00


▼ Exploratory Data Analysis(EDA)

```
df.nunique()


tweet_id          14485
airline_sentiment      3
airline_sentiment_confidence  1023
negativereason        10
negativereason_confidence  1410
airline              6
name                 7701
retweet_count        18
```

```
text                14427
tweet_created       14247
tweet_location      3081
user_timezone       85
dtype: int64
```

```
ax = sns.countplot(x = "airline_sentiment", data = df)
```



```
plt.figure(figsize = (10, 10))
ax = sns.countplot(x = "airline", data = df)
```



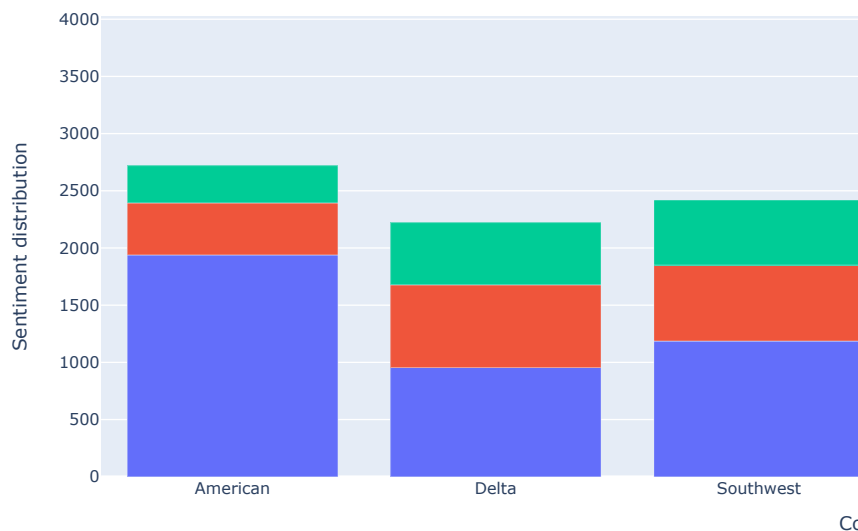
**** Stacked bar chart to show the distribution of reviews per company****

```
import plotly.graph_objects as go
crosstab_sentiments=pd.crosstab(df.airline, df.airline_sentiment)
companies=list(crosstab_sentiments.index)

fig = go.Figure(data=[
    go.Bar(name=col_name, x=companies, y=list(crosstab_sentiments[col_name]))
    for col_name in list(crosstab_sentiments.columns)])
# Change the bar mode
fig.update_layout(barmode='stack',
                  title='Sentiment distribution per company',
                  yaxis=dict(title='Sentiment distribution'),
                  xaxis=dict(title='Companies'))

fig.show()
```

Sentiment distribution per company



Cc

Stacked bar chart to show negative reasons distributions per company

```
crosstab_neg_reasons = pd.crosstab(df["airline"], df["negativereason"])
companies = list(crosstab_neg_reasons.index)

fig = go.Figure(data = [
    go.Bar(name = col_name, x = companies, y = list(crosstab_neg_reasons[col_name]))
    for col_name in list(crosstab_neg_reasons.columns)])

fig.update_layout(barmode = "stack",
                  title = "Negative Reasons Distribution per Company",
                  yaxis = dict(title = "Negative reasons Distribution"),
                  xaxis = dict(title = "Companies"))

fig.show()
```


Negative Reasons Distribution per Company

Pie chart to check the overall distribution for negative reasons.


```

c
labels = list(crosstab_neg_reasons.columns)
values = [crosstab_neg_reasons[col_name].sum() for col_name in labels]

# Use `hole` to create a donut-like pie chart
fig = go.Figure(data=[go.Pie(labels=labels, values=values, hole=.3)])
fig.update_layout(title='Overall distribution for negative reasons')
fig.show()

```

Overall distribution for negative reasons



```
df.drop(df.loc[df["airline_sentiment"] == "neutral"].index, inplace = True)
```

Vectorization Process

```

data = df[
    ["airline_sentiment", "text"]
]
data.head()

```

	airline_sentiment	text	
1	positive	@VirginAmerica plus you've added commercials t...	
3	negative	@VirginAmerica it's really aggressive to blast...	
4	negative	@VirginAmerica and it's a really big bad thing...	
5	negative	@VirginAmerica seriously would pay \$30 a fligh...	
6	positive	@VirginAmerica yes, nearly every time I fly VX...	

```

X = df["text"]
y = df["airline_sentiment"]

```

```

X
1      @VirginAmerica plus you've added commercials t...
3      @VirginAmerica it's really aggressive to blast...
4      @VirginAmerica and it's a really big bad thing...
5      @VirginAmerica seriously would pay $30 a fligh...
6      @VirginAmerica yes, nearly every time I fly VX...
...

```

```

14633 @AmericanAir my flight was Cancelled Flightled...
14634 @AmericanAir right on cue with the delays 🤔
14635 @AmericanAir thank you we got on a different f...
14636 @AmericanAir leaving over 20 minutes Late Flig...
14638 @AmericanAir you have my money, you change my ...
Name: text, Length: 11510, dtype: object

```

y

```

1 positive
3 negative
4 negative
5 negative
6 positive
...
14633 negative
14634 negative
14635 positive
14636 negative
14638 negative
Name: airline_sentiment, Length: 11510, dtype: object

```

Train Test Split

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

```
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(9208,) (2302,) (9208,) (2302,)
```

```
tfidf = TfidfVectorizer(stop_words="english")
```

```
tfidf.fit(X_train)
```

```

▼ TfidfVectorizer
TfidfVectorizer(stop_words='english')

```

```
print(tfidf.get_feature_names_out())
```

```
['00' '000' '000ft' ... 'zv6cfpoh15' 'zvfmxnuelj' 'zzps5ywe2']
```

```
print(tfidf.vocabulary_)
```

```
{'jetblue': 5633, 'did': 3478, 'idea': 5246, 'offered': 7007, 'pay': 7303, 'tix': 9682, 'airport': 1441, 'bc': 1996, 'told': 9717,
```

```
data[data["airline_sentiment"] == "negative"]["text"]
```

```

3 @VirginAmerica it's really aggressive to blast...
4 @VirginAmerica and it's a really big bad thing...
5 @VirginAmerica seriously would pay $30 a fligh...
15 @VirginAmerica SFO-PDX schedule is still MIA.
17 @VirginAmerica I flew from NYC to SFO last we...
...
14631 @AmericanAir thx for nothing on getting us out...
14633 @AmericanAir my flight was Cancelled Flightled...
14634 @AmericanAir right on cue with the delays 🤔
14636 @AmericanAir leaving over 20 minutes Late Flig...
14638 @AmericanAir you have my money, you change my ...
Name: text, Length: 9157, dtype: object

```

```
count_vect = CountVectorizer(stop_words="english")
```

```
neg_matrix = count_vect.fit_transform(data[data["airline_sentiment"]=="negative"]["text"])
```

```
freqs = zip(count_vect.get_feature_names_out(), neg_matrix.sum(axis=0).tolist()[0])
```

```
# Sort from largest to smallest
```

```
print(sorted(freqs, key=lambda x: -x[1])[:100])
```

```
[('flight', 2937), ('united', 2899), ('usairways', 2375), ('americanair', 2089), ('southwestair', 1214), ('jetblue', 1051), ('cance
```

Wordcloud for Positive Reasons

```
new_df = data[data["airline_sentiment"] == "negative"]
words = " ".join(new_df["text"])
cleaned_word = " ".join([word for word in words.split() if "http" not in word and not word.startswith("@") and word != "RT"])
wordcloud = WordCloud(stopwords = STOPWORDS,
                      background_color = "black", width = 3000, height = 2500).generate(cleaned_word)
plt.figure(figsize = (12, 12))
plt.imshow(wordcloud)
plt.axis("off")
plt.show()
```



```
data.drop(data.loc[data["airline_sentiment"] == "neutral"].index, inplace = True)
```



▼ Data Scaling



```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
le.fit(data["airline_sentiment"])
data["airline_sentiment_Encoded"] = le.transform(data["airline_sentiment"])
data.head()
```

	airline_sentiment	text	airline_sentiment_Encoded
1	positive	@VirginAmerica plus you've added commercials t...	1
3	negative	@VirginAmerica it's really aggressive to blast...	0
4	negative	@VirginAmerica and it's a really big bad thing...	0
5	negative	@VirginAmerica seriously would pay \$30 a fligh...	0
6	positive	@VirginAmerica yes, nearly every time I fly VX...	1

```
def tweet_to_words(tweet):
    letters_only = re.sub("[^a-zA-Z]", " ", tweet)
    words = letters_only.lower().split()
    stops = set(stopwords.words("english"))
    meaningful_words = [w for w in words if not w in stops]
    return(" ".join( meaningful_words ))

nltk.download("stopwords")
data["clean tweet"] = data["text"].apply(lambda x: tweet_to_words(x))
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11510 entries, 1 to 14638
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   airline_sentiment                    11510 non-null  object
1   text                                11510 non-null  object
2   airline_sentiment_Encoded            11510 non-null  int64
3   clean_tweet                          11510 non-null  object
dtypes: int64(1), object(3)
memory usage: 449.6+ KB
```

▼ Defining X and y

```
X = data["clean_tweet"]
y = data["airline_sentiment"]

print(X.shape, y.shape)

(11510,) (11510,)

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 42)

print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(8632,) (2878,) (8632,) (2878,)

vect = CountVectorizer()
vect.fit(X_train)

▼ CountVectorizer
CountVectorizer()

X_train_dtm = vect.transform(X_train)
X_test_dtm = vect.transform(X_test)
```

▼ Tuning

```
vect_tunned = CountVectorizer(stop_words = "english", ngram_range = (1, 2), min_df = 0.1, max_df = 0.7, max_features = 100)
vect_tunned

▼ CountVectorizer
CountVectorizer(max_df=0.7, max_features=100, min_df=0.1, ngram_range=(1, 2),
stop_words='english')
```

▼ Model Building

```
from sklearn.svm import SVC
model = SVC(kernel = "linear", random_state = 10)
model.fit(X_train_dtm, y_train)
pred = model.predict(X_test_dtm)

print("Accuracy Score: ", accuracy_score(y_test, pred) * 100)


Accuracy Score: 90.7574704656011

print("Confusion Matrix\n\n", confusion_matrix(y_test, pred))

Confusion Matrix

[[2179 122]
 [ 144 433]]

# Visualizing the confusion matrix using a heatmap
conf_matrix_df = pd.DataFrame(data=conf_matrix, columns=le.classes_, index=le.classes_)
plt.rcParams['figure.figsize'] = [10, 7]
sns.heatmap(conf_matrix_df, annot=True, fmt='d', cmap='Blues', cbar=False, linewidths=0.1, annot_kws={'size': 12})
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
negative	0.94	0.95	0.94	2301
positive	0.78	0.75	0.77	577
accuracy			0.91	2878
macro avg	0.86	0.85	0.85	2878
weighted avg	0.91	0.91	0.91	2878

Conclusions

As we you can see above we have plotted the confusion matrix for predicted sentiments and actual sentiments (negative and positive)

SVM Classifier gives us the best accuracy score i.e 91% precision scores according to the classification report

The confusion matrix shows the TP,TN,FP,FN for sentiments(negative, positive).

▼ Decision Tree Classifier

```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import re
from nltk.corpus import stopwords
from sklearn.preprocessing import LabelEncoder

# Splitting the data into features (X) and target variable (y)
X = data["clean_tweet"]
y = data["airline_sentiment_Encoded"]

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Vectorizing the text data
vect = CountVectorizer()
X_train_dtm = vect.fit_transform(X_train)
X_test_dtm = vect.transform(X_test)

# Initializing and training the Decision Tree model
dt_model = DecisionTreeClassifier(random_state=10)
dt_model.fit(X_train_dtm, y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=10)
```

```
# Predicting on the test set
pred = dt_model.predict(X_test_dtm)
```

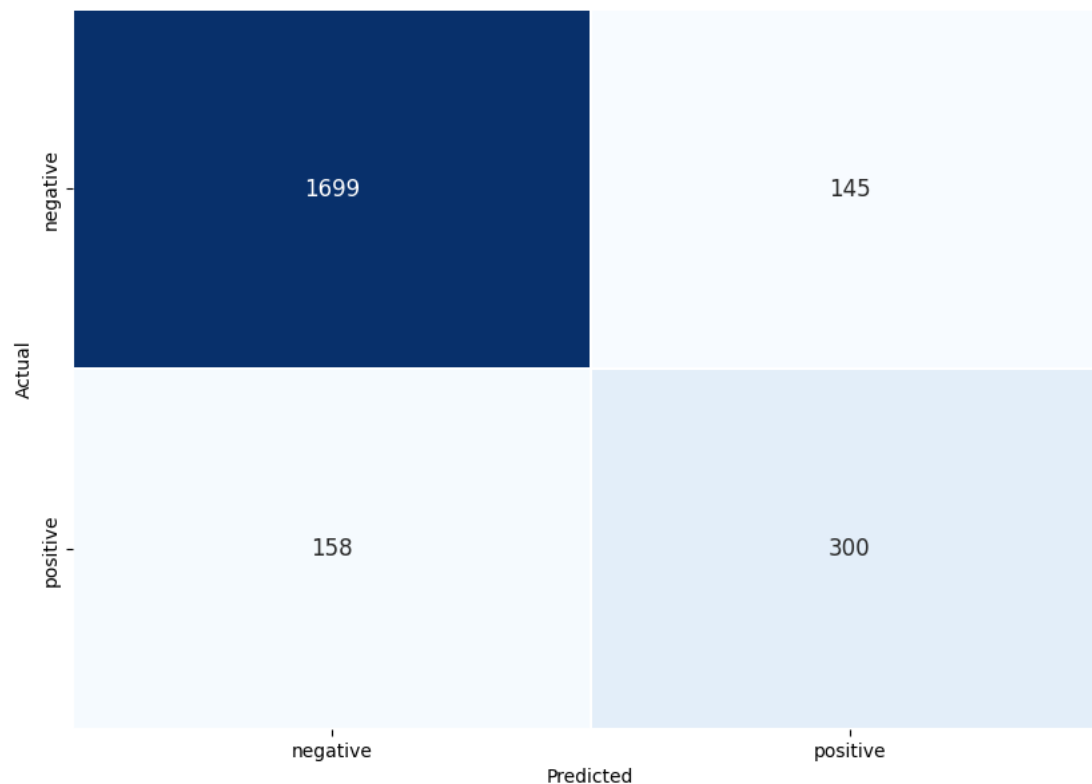
```
# Evaluating the model
accuracy = accuracy_score(y_test, pred)
print("Accuracy Score: {:.2f}%".format(accuracy * 100))
```

Accuracy Score: 86.84%

```
# Displaying the confusion matrix
conf_matrix = confusion_matrix(y_test, pred)
print("Confusion Matrix:\n", conf_matrix)
```

```
Confusion Matrix:
[[1699  145]
 [ 158  300]]
```

```
# Visualizing the confusion matrix using a heatmap
conf_matrix_df = pd.DataFrame(data=conf_matrix, columns=le.classes_, index=le.classes_)
plt.rcParams['figure.figsize'] = [10, 7]
sns.heatmap(conf_matrix_df, annot=True, fmt='d', cmap='Blues', cbar=False, linewidths=0.1, annot_kws={'size': 12})
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Displaying the classification report
report = classification_report(y_test, pred, target_names=le.classes_)
print("Classification Report:\n", report)
```

```
Classification Report:
              precision    recall  f1-score   support

 negative      0.91      0.92      0.92     1844
 positive      0.67      0.66      0.66      458

 accuracy              0.87              0.87     2302
 macro avg              0.79              0.79              0.79     2302
 weighted avg              0.87              0.87              0.87     2302
```

Conclusion:

The Decision Tree Classifier exhibits a reasonable overall performance with an accuracy of 87%.** This suggests that the model effectively predicts the correct class for a majority of instances based on the given features.**

▼ Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd


# Initialize Random Forest model
model = RandomForestClassifier(random_state=10)
model.fit(X_train_dtm, y_train)
pred = model.predict(X_test_dtm)

# Calculate and print accuracy
accuracy = accuracy_score(y_test, pred)
print("Accuracy Score: {:.2f}%".format(accuracy * 100))

# Display confusion matrix
conf_matrix = confusion_matrix(y_test, pred)
print("Confusion Matrix:\n", conf_matrix)

Confusion Matrix:
[[2220  81]
 [ 204 373]]

# Visualize confusion matrix using a heatmap
conf_matrix_df = pd.DataFrame(data=conf_matrix, columns=model.classes_, index=model.classes_)
plt.rcParams['figure.figsize'] = [15, 8]
sns.heatmap(conf_matrix_df, annot=True, fmt='d', cmap='Paired', cbar=False, linewidths=0.1, annot_kws={'size': 25})
plt.xticks(fontsize=20)
plt.yticks(fontsize=20)
plt.show()
```



```
# Display classification report
report = classification_report(y_test, pred)
print("Classification Report:\n", report)
```


Classification Report:				
	precision	recall	f1-score	support
negative	0.92	0.96	0.94	2301
positive	0.82	0.65	0.72	577
accuracy			0.90	2878
macro avg	0.87	0.81	0.83	2878
weighted avg	0.90	0.90	0.90	2878

Conclusion:

The Random Forest Classifier demonstrates strong overall accuracy, achieving a score of 90%. This indicates the model's ability to correctly classify instances into their respective classes based on the given features.

While the model performs well in identifying the "negative" class with high precision and recall (0.92 and 0.96, respectively), it faces challenges in accurately classifying the "positive" class, showing lower precision (0.82) and recall (0.65). This suggests potential areas for improvement, particularly in capturing positive instances.

▼ K-Nearest-Neighbors(KNN)

```

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import re
from nltk.corpus import stopwords
from sklearn.preprocessing import LabelEncoder

# Splitting the data into features (X) and target variable (y)
X = data["clean_tweet"]
y = data["airline_sentiment_Encoded"]

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Vectorizing the text data
vect = CountVectorizer()
X_train_dtm = vect.fit_transform(X_train)
X_test_dtm = vect.transform(X_test)

# Initializing and training the KNN model
knn_model = KNeighborsClassifier()
knn_model.fit(X_train_dtm, y_train)

▼ KNeighborsClassifier
KNeighborsClassifier()

# Predicting on the test set
pred = knn_model.predict(X_test_dtm)

# Evaluating the model
accuracy = accuracy_score(y_test, pred)
print("Accuracy Score: {:.2f}%".format(accuracy * 100))

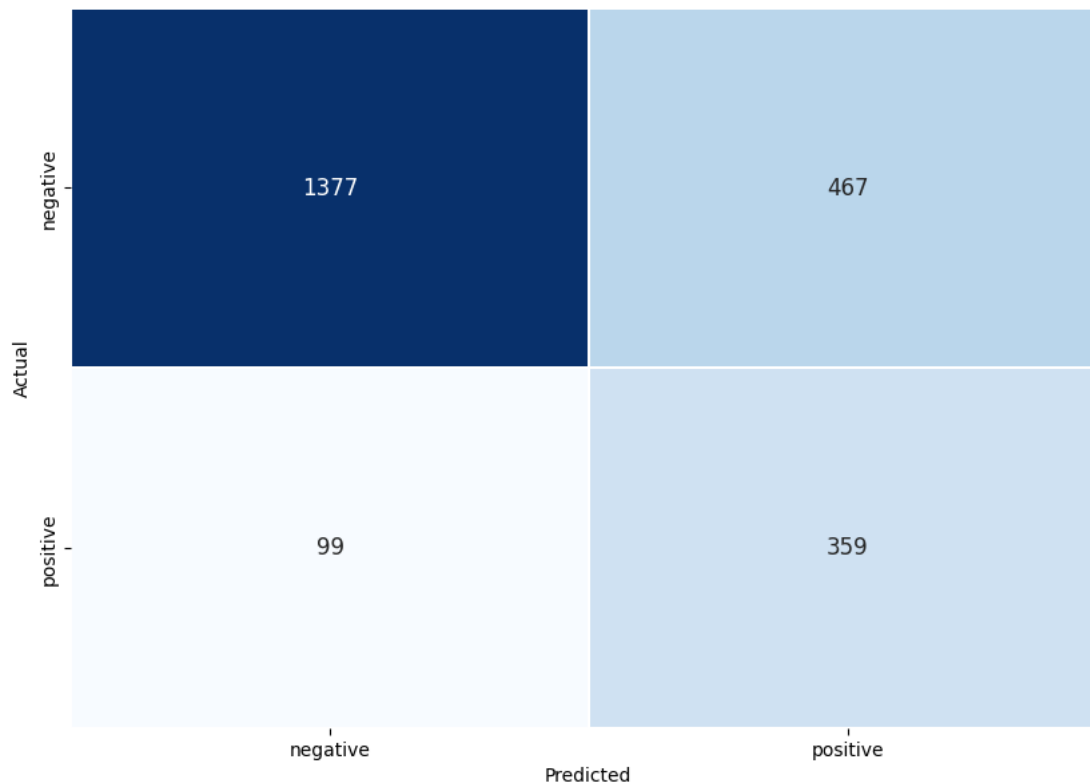
Accuracy Score: 75.41%

# Displaying the confusion matrix
conf_matrix = confusion_matrix(y_test, pred)
print("Confusion Matrix:\n", conf_matrix)

Confusion Matrix:
[[1377  467]
 [  99 359]]

```

```
# Visualizing the confusion matrix using a heatmap
conf_matrix_df = pd.DataFrame(data=conf_matrix, columns=le.classes_, index=le.classes_)
plt.rcParams['figure.figsize'] = [10, 7]
sns.heatmap(conf_matrix_df, annot=True, fmt='d', cmap='Blues', cbar=False, linewidths=0.1, annot_kws={'size': 12})
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Displaying the classification report
report = classification_report(y_test, pred, target_names=le.classes_)
print("Classification Report:\n", report)
```

```
Classification Report:
              precision    recall  f1-score   support

 negative      0.93      0.75      0.83     1844
 positive      0.43      0.78      0.56      458

 accuracy              0.75     2302
 macro avg      0.68      0.77      0.69     2302
 weighted avg    0.83      0.75      0.78     2302
```

▼ Naive Bayes Classifier

```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import re
from nltk.corpus import stopwords
from sklearn.preprocessing import LabelEncoder

# Splitting the data into features (X) and target variable (y)
X = data["clean_tweet"]
y = data["airline_sentiment_Encoded"]

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Vectorizing the text data
vect = CountVectorizer()
X_train_dtm = vect.fit_transform(X_train)
X_test_dtm = vect.transform(X_test)
```

```
# Initializing and training the Naive Bayes model
nb_model = MultinomialNB()
nb_model.fit(X_train_dtm, y_train)
```

```
▸ MultinomialNB
MultinomialNB()
```

```
# Predicting on the test set
pred_nb = nb_model.predict(X_test_dtm)
```


```
# Evaluating the model
accuracy_nb = accuracy_score(y_test, pred_nb)
print("Naive Bayes - Accuracy Score: {:.2f}%".format(accuracy_nb * 100))
```

```
Naive Bayes - Accuracy Score: 91.88%
```

```
# Displaying the confusion matrix
conf_matrix_nb = confusion_matrix(y_test, pred_nb)
print("Naive Bayes - Confusion Matrix:\n", conf_matrix_nb)
```

```
Naive Bayes - Confusion Matrix:
[[1810  34]
 [ 153 305]]
```

```
# Visualizing the confusion matrix using a heatmap
conf_matrix_df_nb = pd.DataFrame(data=conf_matrix_nb, columns=le.classes_, index=le.classes_)
plt.rcParams['figure.figsize'] = [10, 7]
sns.heatmap(conf_matrix_df_nb, annot=True, fmt='d', cmap='Blues', cbar=False, linewidths=0.1, annot_kws={'size': 12})
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
# Displaying the classification report
report_nb = classification_report(y_test, pred_nb, target_names=le.classes_)
print("Naive Bayes - Classification Report:\n", report_nb)
```

```
Naive Bayes - Classification Report:
              precision    recall  f1-score   support

negative     0.92         0.98         0.95         1844
```

positive	0.90	0.67	0.77	458
accuracy			0.92	2302
macro avg	0.91	0.82	0.86	2302
weighted avg	0.92	0.92	0.91	2302

Conclusion:

The **Naive Bayes** Classifier demonstrates strong overall accuracy as compared to other ML models, achieving an accuracy score of **92%**. This indicates the model's ability to correctly classify instances into their respective classes based on the given features.

The classifier performs exceptionally well in identifying the "negative" class, with high precision (0.92) and recall (0.98). However, there is room for improvement in classifying the "positive" class, as reflected by a slightly lower precision (0.90) and recall (0.67). This suggests potential areas for fine-tuning, particularly in capturing positive instances.

The macro and weighted average F1-scores are both 0.86, indicating a balanced performance in terms of precision and recall across classes. The model provides a robust trade-off between precision and recall, offering a comprehensive evaluation of its effectiveness.